



CUDA Workshop

Nuno Subtil
nsubtil@nvidia.com



Getting Started with CUDA



- **Requirements:**
 - **CUDA-capable GPU**
 - **C compiler**
 - **Windows: Microsoft Visual Studio**
(Express Edition is fine --- but mind the license!)
 - **Linux: gcc 4.x**
 - **Mac OS X: Xcode 3.2.x**
 - **CUDA drivers and toolkit**
 - **Your favorite text editor**

IDEs for CUDA



- Pretty much whatever C IDE you are comfortable with
- Compile with `nvcc` instead of the native compiler
 - `.c`, `.cpp`, `.cu` files should be handled correctly without intervention
- Windows: Parallel Nsight
 - Integrated build and debug environment for Visual Studio
 - Not really a requirement

CUDA Basics: Driver vs. Runtime API



- **Driver API**
 - Pure low-level C API (much like OpenGL)
 - No syntax extensions
 - Verbose and not really very friendly
- **Runtime API**
 - Implemented on top of driver API
 - Extensions to C syntax (think OpenMP)
 - A lot more friendly with no real loss of flexibility

First Step: Hello World!



```
__global__ void empty_kernel(void) {  
}
```

```
int main( void ) {  
    empty_kernel<<<1, 1>>>>();  
    printf("Hello World!\n");  
    return 0;  
}
```

```
nvcc -o hello hello.cu
```

GPU Kernel



```
__global__ void empty_kernel(void) {  
}
```

- **__global__** means “will run on GPU when called from CPU”
- Compiler automatically generates GPU code for the function
- Also available: **__device__** (function can be called on GPU but not on CPU)

Kernel Launch Syntax

```
empty_kernel<<<1, 1>>>();
```

- Automatically invokes GPU code
- Kernels run on a “grid”
- Launch syntax specifies the grid
- Function parameters work just like in C

Second Step: 2 + 7




```
__global__ void add(int a, int b, int *c) {  
    *c = a + b;  
}  
  
int main( void ) {  
    int c, *dev_c;  
    HANDLE_ERROR( cudaMalloc(&dev_c, sizeof(int)) );  
    add<<<1, 1>>>(2, 7, dev_c);  
    HANDLE_ERROR( cudaMemcpy(&c, dev_c, sizeof(int),  
                             cudaMemcpyDeviceToHost) );  
    printf("2 + 7 = %d\n", c);  
}
```



Host vs Device Memory



- **Device:** your graphics/compute card
- **Host:** the rest of your computer




Host vs Device Memory



- **NUMA: separate host and device address spaces**
 - Device can not read/write host memory (well... sort of)
 - Host can not read/write device memory (this one is mostly true)
- **Full 32-bit device pointers**
- **Device-specific memory management functions**

Device Pointers



- **NOT** valid on host!
 - ... and host pointers **not** valid on device either!
- Pointer arithmetic is allowed
- No bounds checking
 - You can segfault on the GPU and your process can get killed for that!)

Become a CUDA Memory Management Expert



- **cudaMalloc: allocate on-device memory**
 - Runs on host, returns device pointer to host
- **cudaMemcpy:**
 - Copy between host and device
 - Copy between two device buffers
 - Copy between different devices (on a multi-GPU system)
 - Called on host, executes on host and on device
 - (Also available: cudaMemcpySet)
- **cudaFree: free device memory**

2 + 7 = ?...



```
int main( void ) {  
    int c, *dev_c;  
    HANDLE_ERROR( cudaMalloc(&dev_c, sizeof(int)) );  
    add<<<1, 1>>>(2, 7, dev_c);  
    HANDLE_ERROR( cudaMemcpy(&c, dev_c, sizeof(int),  
                             cudaMemcpyDeviceToHost) );  
    printf("2 + 7 = %d\n", c);  
}
```

2 + 7 = 27...



```
int main( void ) {  
    int c, *dev_c;  
    HANDLE_ERROR( cudaMalloc(&dev_c, sizeof(int)) );  
    add<<<1, 1>>>(2, 7, dev_c);  
    HANDLE_ERROR( cudaMemcpy(&c, dev_c, sizeof(int),  
                             cudaMemcpyDeviceToHost) );  
    printf("2 + 7 = %d\n", c);  
}
```

Host Memory

Device Pointer

2 + 7 = 9!




```
int main( void ) {  
    int c, *dev_c;  
    HANDLE_ERROR( cudaMalloc(&dev_c, sizeof(int)) );  
    add<<<1, 1>>>(2, 7, dev_c);  
    HANDLE_ERROR( cudaMemcpy(&c, dev_c, sizeof(int),  
                             cudaMemcpyDeviceToHost) );  
  
    cudaFree(dev_c);  
    printf("2 + 7 = %d\n", c);  
}
```

Vector Sum




Problem: compute the sum of two integer vectors



Serial CPU Version



```
void sum(int *a,  
         int *b, int *c)  
{  
    for(int i = 0; i < N; i++) {  
        c[i] = a[i] + b[i];  
    }  
}
```




Serial Approach on the GPU



Exercise: run our vector sum on the GPU using a single thread


- **Implement single-threaded GPU kernel**
 - Hint: very similar to CPU version
- **Allocate and fill GPU memory**
- **Launch kernel on the GPU**
- **Read results back**



Serial Approach on the GPU

To run on GPU, just add `__global__`...

```
__global__ void sum(int *a,  
                    int *b, int *c)  
{  
    for(i = 0; i < N; i++) {  
        c[i] = a[i] + b[i];  
    }  
}
```



Serial Approach on the GPU



... do the memory allocation on the GPU and copy the data ...

```
int *dev_a, *dev_b, *dev_c;  
HANDLE_ERROR( cudaMalloc(&dev_a, sizeof(int) * N) );  
HANDLE_ERROR( cudaMalloc(&dev_b, sizeof(int) * N) );  
HANDLE_ERROR( cudaMalloc(&dev_c, sizeof(int) * N) );  
cudaMemcpy(dev_a, host_a, sizeof(int) * N,  
           cudaMemcpyHostToDevice);  
cudaMemcpy(dev_b, host_b, sizeof(int) * N,  
           cudaMemcpyHostToDevice);
```


Serial Approach on the GPU

... then launch the kernel...

```
sum<<<1, 1>>>(dev_a, dev_b, dev_c);
```

... and finally read the results back.

```
cudaMemcpy(host_c, dev_c, sizeof(int) * N,  
           cudaMemcpyDeviceToHost);
```



Running in Parallel



Going Parallel




- **CPUs: single-threaded, complex code**
 - Generally one, sometimes two threads per CPU core
 - Hardware copes very well with “bad” code
- **GPUs: massively multi-threaded, simple code**
 - Not that great at branching (but getting better!)
 - Run thousands of threads per GPU core very efficiently
 - Brute force makes up for the limitations 😊

Parallel Approach on the GPU



- Run one thread per index
- Each thread computes a single element of the output




Parallel Approach on the GPU



Kernel launch syntax:

```
sum<<<N, 1>>>>(dev_a, dev_b,  
dev_c) ;
```



Parallel Approach on the GPU




Kernel launch syntax:

```
sum<<<N, 1>>>>(dev_a, dev_b,  
dev_c) ;
```

Blocks

Threads per block



Parallel Approach on the GPU




Kernel launch syntax:

```
sum<<<N, 1>>>(dev_a, dev_b,  
dev_c) ;
```

Blocks

Threads per block



One thread per block


Parallel Approach on the GPU



Kernel function runs once per element

- How do we know which element to process?


A: CUDA tells us through `blockIdx`



Parallel Approach on the GPU



```
__global__ void sum(int *a,  
                    int *b, int *c)  
{  
    /* magic here */  
}
```




Exercise: implement the parallel version of the kernel

Parallel Approach on the GPU



```
__global__ void sum(int *a,  
                    int *b, int *c)  
{  
    int i = blockIdx.x;  
    c[i] = a[i] + b[i];  
}
```



Q: How big can the vector be?

GPU Work Partitioning



GPU Work Partitioning



- Each GPU kernel runs in thousands of threads in parallel
- Sets of threads (“thread blocks”) will share certain hardware resources
- CUDA runtime mimics the hardware architecture
 - Kernels run by blocks of threads
 - Launch syntax specifies both

GPU Work Partitioning




Kernel launch syntax:

```
sum<<<N, 1>>>(dev_a, dev_b,  
dev_c) ;
```

Blocks

Threads per block



One thread per block


GPU Work Partitioning



Kernel launch syntax:

```
sum<<<N/2, 2>>>(dev_a, dev_b,  
dev_c);
```

2 threads/block * N/2 blocks
= N threads




GPU Work Partitioning



Kernel launch syntax:

```
sum<<<N/4, 4>>>(dev_a, dev_b,  
dev_c);
```

4 threads/block * N/4 blocks
= N threads



GPU Work Partitioning: Why?



- **Immediate concern: limited number of blocks per launch**
 - Vector sum example breaks!
- **Efficiency**
 - Maximize hardware occupancy
 - Make use of shared memory per block
 - Benefit from memory access coalescing
 - Lower instruction fetch pressure


GPU Work Partitioning



Kernel function still runs once per element

- How do we know which element to process?

A: `blockIdx`, `blockDim`, `threadIdx`.



GPU Work Partitioning




```
sum<<<N/2, 2>>>(dev_a, dev_b,  
dev_c);
```

blockIdx.x

threadIdx.x

Vector element



GPU Work Partitioning




blockIdx: block index
threadIdx: thread index within block
blockDim: threads per block (2)
gridDim: blocks per launch (N/2)

blockIdx.x

threadIdx.x

Vector element



GPU Work Partitioning




blockIdx.x

threadIdx.x

Vector element

```
__global__ void sum(int *a,  
                    int *b, int *c)  
{  
    /* ??? */  
}
```


Exercise: make it work!



GPU Work Partitioning



```
__global__ void sum(int *a,  
                    int *b, int *c)  
{  
    int i;  
    i = blockIdx.x * blockDim.x +  
        threadIdx.x;  
    c[i] = a[i] + b[i];  
}
```



Kernel Launch Limits



How many threads can you launch at once now?

A: Many, but not infinitely many!

Kernel Launch Limits



- **Limited number of threads per block**
 - GF100: maximum is 1024 threads
 - Depends on shared resource usage
- **Limited number of blocks per launch**
 - GF100: maximum is 65536 blocks
- **Total number of threads is still limited**
- **CUDA provides an API to query these limits**

Querying Device Limits

```
cudaDeviceProp prop;  
HANDLE_ERROR( cudaDeviceGetProperties(&prop, 0) );
```

Interesting fields:

- `prop.name` (device name string)
- `prop.totalGlobalMem` (size of device memory)
- `prop.maxThreadsPerBlock`
- `prop.gridDim[0]` (max blocks per launch)

See the `deviceQuery` example from the CUDA SDK

Arbitrarily long vector sums



Modify the vector sum example to handle arbitrary size vectors while adhering to your particular device's limits.

(Hint: you can do more work per thread. But you don't have to.)

Arbitrarily long vector sums

Run the kernel multiple times...

```
__global__ void sum(int *a, int *b, int *c, int o) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x + o;  
    if (i < N)  
        c[i] = a[i] + b[i];  
}
```

...

```
for(j = 0; j < N; j += 256 * 256)  
    sum<<<256, 256>>>(dev_a, dev_b, dev_c, j);
```

Arbitrarily long vector sums

... or process more elements per thread (how does this work?)

```
__global__ void sum(int *a, int *b, int *c) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    while (i < N) {  
        c[i] = a[i] + b[i];  
        i += blockDim.x * blockDim.x;  
    }  
}
```

...

```
sum<<<B, T>>>>(dev_a, dev_b, dev_c);
```


GPU Reductions



Vector Dot Product




Given two arbitrarily long integer vectors, compute their dot product.

Vector Dot Product




$$A[0] * B[0] + A[1] * B[1] + .. + A[N-1] * B[N-1]$$



Parallel Vector Dot Product



- **Two steps**
 - Multiplication
 - Summation
- **Summation is a “reduction” operation**
 - Many input elements for one output element



**Exercise: implement dot product on the GPU.
(Do the reduction step on the CPU for now.)**

Parallel Vector Dot Product: Reduction



- Obtain one output element from many input elements
 - Input: a vector
 - Output: a scalar (sum of elements of the vector)
- Solution: use many threads and multiple steps

Reduction




Input:



Reduction

Input:

Reduction step 1:




Reduction

Input:

Reduction step 1:

Reduction step 2:




Reduction



Input:

Reduction step 1:

Reduction step 2:



... but we must not start
step 2 until all threads are
done with step 1!

Thread Cooperation: Barrier

CUDA implements a thread barrier: `__syncthreads()`

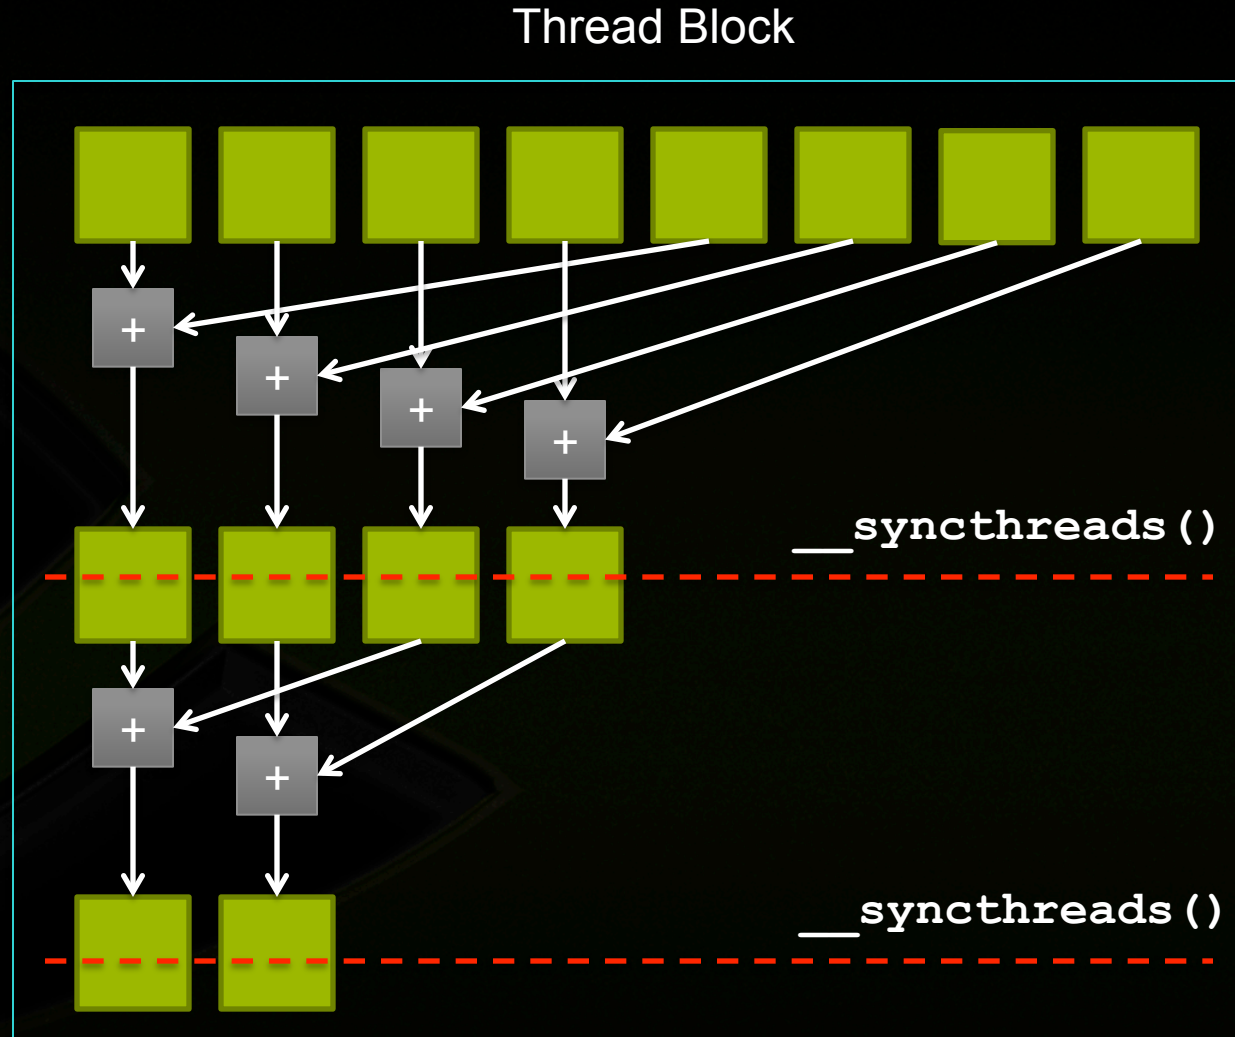
- Threads wait inside `__syncthreads()` until *all* threads *within the same block* reach `__syncthreads()`
 - If one thread in the block calls `__syncthreads()`, then ALL threads on the same block must call `__syncthreads()`
- No synchronization across blocks
 - CUDA requires that blocks be completely independent
 - Order of execution of blocks is arbitrary --- not guaranteed to be scheduled in order

Reduction


Input:

Reduction step 1:

Reduction step 2:



Reduction



Exercise: implement GPU reduction for the dot product.


GPU Memory Hierarchy and Thread Cooperation



GPU Memory Hierarchy




- **Each thread has local storage**
 - On-chip, low-latency, very fast
 - Implemented as registers with spill over into L1 cache
 - GPU function variables live here



GPU Memory Hierarchy




- Thread blocks have shared memory
 - On-chip, fast
 - Implemented as L2 cache
 - GPU variables declared as `__shared__` live here



GPU Memory Hierarchy



- GPU has access to global memory
 - RAM on the graphics card
 - Off-chip, much slower
 - `cudaMalloc` allocates memory here



Thread Cooperation via Shared Memory


Shared memory: common storage area visible to all threads within a block

- **Just declare a variable or array with the `__shared__` attribute...**

```
__shared__ int localStorage[THREADS_PER_BLOCK];
```

- **... and then make sure to use it wisely!**
 - Limited amount of shared memory available

Dot Product with Shared Memory



Exercise: make use of shared memory instead of global memory for the reduction in the dot product.

Atomic Operations



Thread Cooperation: Atomics



- **__syncthreads()** is not useful for synchronizing across thread blocks
- **CUDA provides atomic memory access functions**
 - Read-modify-write operations on memory addresses
 - Can operate on global or shared memory

Thread Cooperation: Atomics

```
int atomicAdd(int *address, int val)
```

Reads the old value at address

Computes `val + old` and writes result back to address

Returns the old value

```
int atomicSub(int *address, int val)
```

Same as `atomicAdd`, but performs a subtraction

Thread Cooperation: Atomics

```
int atomicExch(int *address, int val)
```

Writes val at address, returns old value

```
int atomicCAS(int *address, int cmp, int val)
```

Reads old value at address

```
result = (old == cmp ? val : old)
```

Writes result to address, returns old value

Also available: atomicInc, atomicDec, atomicAnd/Or/Xor

Thread Cooperation: Atomics



- **When do we use atomics?**
 - Synchronizing access to a memory area
 - Allows multiple threads to communicate safely
- **When not to use atomics? When you need performance!**
 - Atomics serialize threads
 - Reduced parallelism
 - Performance penalty is especially high when using atomics on global memory

Simple Atomics Example: Vector Sum



```
__global__ void vector_sum(int *vector, int *out) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    int sum = 0;  
    while(i < N) {  
        sum += vector[i];  
        i += blockDim.x * gridSize.x;  
    }  
    atomicAdd(out, sum);  
}
```

Histogram Computation



Problem: compute the histogram of an array of bytes.

(In other words, count how many bytes in the given data set have each of the 256 possible values.)

Histogram: CPU Solution

```
void compute_histogram(char *data, int *histogram) {  
    int c;  
    memset(histogram, 0, sizeof(int) * 256);  
    for(c = 0; c < N; c++)  
        histogram[data[c]]++;  
}
```

Now implement the GPU version. And make it fast!

(Hint: shared memory is your friend)